
DIFFERENTIABLE SAMPLING OF CATEGORICAL DISTRIBUTIONS USING THE CATLOG-DERIVATIVE TRICK

Lennert De Smet
KU Leuven

Emanuele Sansone
KU Leuven

Pedro Zuidberg Dos Martires
Örebro University

ABSTRACT

Categorical random variables can faithfully represent the discrete and uncertain aspects of data as part of a discrete latent variable model. Learning in such models necessitates taking gradients with respect to the parameters of the categorical probability distributions, which is often intractable due to their combinatorial nature. A popular technique to estimate these otherwise intractable gradients is the Log-Derivative trick. This trick forms the basis of the well-known REINFORCE gradient estimator and its many extensions. While the Log-Derivative trick allows us to differentiate through samples drawn from categorical distributions, it does not take into account the discrete nature of the distribution itself. Our first contribution addresses this shortcoming by introducing the CatLog-Derivative trick – a variation of the Log-Derivative trick tailored towards categorical distributions. Secondly, we use the CatLog-Derivative trick to introduce IndeCateR, a novel and unbiased gradient estimator for the important case of products of independent categorical distributions with provably lower variance than REINFORCE. Thirdly, we empirically show that IndeCateR can be efficiently implemented and that its gradient estimates have significantly lower bias and variance for the same number of samples compared to the state of the art.

1 Introduction

Categorical random variables naturally emerge in many domains in AI, such as language modelling, reinforcement learning and neural-symbolic AI Garcez et al. [2022]. They are compelling because they can faithfully represent the discrete concepts present in data in a sound probabilistic fashion. Unfortunately, inference in probabilistic models with categorical latent variables is usually computationally intractable due to its combinatorial nature. This intractability often leads to the use of sampling-based, approximate inference techniques, which in turn poses problems to gradient-based learning as sampling is an inherently non-differentiable process.

In order to bypass this non-differentiability, two main classes of gradient estimators have been developed. On the one hand, there is a range of unbiased estimators based on the Log-Derivative trick and the subsequent REINFORCE gradient estimator Williams [1992]. On the other hand, we have biased estimators that use continuous relaxations to which the reparametrisation trick Ruiz et al. [2016] can be applied, such as the Gumbel-Softmax trick Jang et al. [2017], Maddison et al. [2017].

A clear advantage of the REINFORCE estimator over relaxation-based estimators is its unbiased nature. However, REINFORCE tends to be sample-inefficient and its gradient estimates exhibit high variance in practice. To resolve these issues, methods have been proposed that modify REINFORCE by, for instance, adding control variates Richter et al. [2020], Titsias, Michalis and Shi, Jiaxin [2022]. These modified estimators have been shown to deliver more robust gradient estimates than standard REINFORCE.

Instead of modifying REINFORCE, we take a different approach and modify the Log-Derivative trick by explicitly taking into account that we are working with multivariate categorical distributions. We call this first contribution the *CatLog-Derivative trick*. Interestingly, we show that our CatLog-Derivative trick leads to Rao-Blackwellised estimators Casella and Robert [1996], immediately giving us a guaranteed reduction in variance. The CatLog-Derivative trick can also be seen as a generalisation of the Local Expectation Gradients (LEG) Titsias, Michalis and Lázaro-Gredilla, Miguel [2015] capable of exploiting more structural, distributional properties. This connection to LEG will be clarified throughout the paper.

As a second contribution, we propose IndeCateR (read as ‘indicator’), a gradient estimator for the special case of independent categorical random variables. IndeCateR is a hyperparameter-free estimator that can be implemented efficiently by leveraging parallelisation on modern graphical processing units (GPUs). Thirdly, we empirically show that IndeCateR is competitive with comparable state-of-the-art gradient estimators on a range of standard benchmarks from the literature.

2 Notation and Preliminaries

Throughout this paper, we consider expectations with respect to multivariate categorical probability distributions of the form

$$\mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})] = \sum_{\mathbf{x} \in \Omega(\mathbf{X})} p(\mathbf{x}) f(\mathbf{x}), \quad (2.1)$$

where we assume this expectation to be finite. The symbol \mathbf{X} denotes a random vector (X_1, \dots, X_D) of D categorical random variables while $p(\mathbf{X})$ denotes a multivariate probability distribution. The expression $\mathbf{X} \sim p(\mathbf{X})$ indicates that the random vector \mathbf{X} is distributed according to $p(\mathbf{X})$. On the right-hand side of Equation (2.1) we write the expectation as an explicit sum over $\Omega(\mathbf{X})$, the finite sample space of the random vector \mathbf{X} , using $\mathbf{x} = (x_1, \dots, x_D)$ for the specific assignments of the random vector (X_1, \dots, X_D) .

Given an order of the random variables in \mathbf{X} , we can induce a factorisation of the joint probability distribution as follows

$$p(\mathbf{X}) = \prod_{d=1}^D p(X_d \mid \mathbf{X}_{< d}). \quad (2.2)$$

Here, $\mathbf{X}_{< d}$ denotes the ordered set of random variables (X_1, \dots, X_{d-1}) . Similarly, $\mathbf{X}_{> d}$ will denote the ordered set (X_{d+1}, \dots, X_D) in subsequent sections.

When performing gradient-based learning, we are interested in partial derivatives of the expected value in (2.1), i.e., $\partial_\lambda \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})]$. Here, we take the partial derivative of the expectation with respect to the parameter λ and assume that the distribution $p(\mathbf{X})$ and the function $f(\mathbf{X})$ depend on a set of parameters Λ with $\lambda \in \Lambda$. For probability distributions to which the reparametrisation trick does not apply, we can rewrite the partial derivative using the Log-Derivative trick.

Theorem 2.1 (Log-Derivative Trick Williams [1992]). Let $p(\mathbf{X})$ be a probability distribution and $f(\mathbf{X})$ such that its expectation is finite, with both functions depending on a set of parameters Λ . Then, it holds that

$$\partial_\lambda \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})] = \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [\partial_\lambda f(\mathbf{X})] + \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X}) \partial_\lambda \log p(\mathbf{X})]. \quad (2.3)$$

In general, both expectations in Equation (2.3) are intractable and often estimated with a Monte Carlo scheme. The most immediate such estimation is provided by the REINFORCE gradient estimator Williams [1992]

$$\partial_\lambda \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})] \approx \frac{1}{N} \sum_{n=1}^N \left(\partial_\lambda f(\mathbf{x}^{(n)}) + f(\mathbf{x}^{(n)}) \partial_\lambda \log p(\mathbf{x}^{(n)}) \right). \quad (2.4)$$

The superscript on $\mathbf{x}^{(n)}$ denotes that it is the n^{th} sample vector drawn from $p(\mathbf{X})$.

A well-known problem with the REINFORCE gradient estimator is the high variance stemming from the second term in Equation (2.4). A growing body of research has been tackling this problem by proposing variance reduction techniques Grathwohl et al. [2018], Richter et al. [2020], Titsias, Michalis and Shi, Jiaxin [2022], Tucker et al. [2017]. In what follows we will focus on estimating this second term and drop the first term, since it can be assumed to be unproblematic.

3 The CatLog-Derivative Trick

The standard Log-Derivative trick and its corresponding gradient estimators are applicable to both discrete and continuous probability distributions. However, this generality limits their usefulness when it comes to purely categorical random variables. For example, the REINFORCE gradient estimator suffers from high variance when applied to problems involving high-dimensional multivariate categorical random variables. In such a setting there are exponentially many possible states to be sampled, which makes it increasingly unlikely that a specific state gets sampled. We now introduce the CatLog-Derivative trick that reduces the exponential number of states arising in a multivariate categorical distribution by exploiting the distribution’s factorisation.

Theorem 3.1 (CatLog-Derivative Trick). Let $p(\mathbf{X})$ be a multivariate categorical probability distribution that depends on a set of parameters Λ and assume $\mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})]$ is finite. Then, it holds that $\partial_\lambda \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})]$ is equal to

$$\sum_{d=1}^D \sum_{x_\delta \in \Omega(X_d)} \mathbb{E}_{\mathbf{X}_{< d} \sim p(\mathbf{X}_{< d})} [\partial_\lambda p(x_\delta | \mathbf{X}_{< d}) \mathbb{E}_{\mathbf{X}_{> d} \sim p(\mathbf{X}_{> d} | x_\delta, \mathbf{X}_{< d})} [f(\mathbf{X}_{\neq d}, x_\delta)]] \quad (3.1)$$

Proof. We start by applying the standard Log-Derivative trick and fill in the product form of the categorical distribution followed by pulling this product out of the logarithm and the expectation

$$\mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X}) \partial_\lambda \log p(\mathbf{X})] = \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} \left[f(\mathbf{X}) \partial_\lambda \log \prod_{d=1}^D p(X_d | \mathbf{X}_{< d}) \right], \quad (3.2)$$

$$= \sum_{d=1}^D \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X}) \partial_\lambda \log p(X_d | \mathbf{X}_{< d})]. \quad (3.3)$$

To continue, we write out the expectation explicitly and write the sum for the random variable X_d separately, resulting in

$$\sum_{d=1}^D \sum_{\mathbf{x} \in \Omega(\mathbf{X}_{\neq d})} \sum_{x_\delta \in \Omega(X_d)} p(\mathbf{x}_{\neq d}, x_\delta) f(\mathbf{x}_{\neq d}, x_\delta) \partial_\lambda \log p(x_\delta | \mathbf{x}_{< d}). \quad (3.4)$$

Next, we factorize the joint probability $p(\mathbf{x}_{\neq d}, x_\delta)$ as $p(\mathbf{x}_{> d} | x_\delta, \mathbf{x}_{< d}) p(x_\delta | \mathbf{x}_{< d}) p(\mathbf{x}_{< d})$. Multiplying the second of these factors with $\partial_\lambda \log p(x_\delta | \mathbf{x}_{< d})$ gives us $\partial_\lambda p(x_\delta | \mathbf{x}_{< d})$. Finally, plugging $\partial_\lambda p(x_\delta | \mathbf{x}_{< d})$ into Equation (3.4) gives the desired expression. \square

Intuitively, the CatLog-Derivative trick decomposes the Log-Derivative trick into an explicit sum of multiple Log-Derivative tricks, one for each of the categorical random variables present in the multivariate distribution. Next, we use the CatLog-Derivative trick to define a novel gradient estimator that exploits the structure of a probability distribution by following the variable ordering of the distribution’s factorisation and effectively Rao-Blackwellises Casella and Robert [1996] REINFORCE.

Definition 3.2 (The SCatR gradient estimator). We define the Structured Categorical REINFORCE (SCatR) estimator via the expression

$$\partial_\lambda \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})] \approx \sum_{d=1}^D \sum_{x_\delta \in \Omega(X_d)} \frac{1}{N} \sum_{n_d=1}^N \partial_\lambda p(x_\delta | \mathbf{x}_{< d}^{(n_d)}) f(\mathbf{x}_{< d}^{(n_d)}, x_\delta, \mathbf{x}_{> d}^{(n_d)}), \quad (3.5)$$

where the sample $\mathbf{x}_{> d}^{(n_d)}$ is drawn while conditioning on x_δ and $\mathbf{x}_{< d}^{(n_d)}$. The subscript on n_d indicates that *different samples* can be drawn for every d .

Proposition 3.3. The SCatR estimator Rao-Blackwellises REINFORCE.

Proof. We start from Equation (2.4) for the REINFORCE estimator, where we ignore the first term and factorize the probability distribution similar to Equation (3.3)

$$\partial_\lambda \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})] \approx \frac{1}{N} \sum_{n=1}^N f(\mathbf{x}^{(n)}) \partial_\lambda \log p(\mathbf{x}^{(n)}) \approx \sum_{d=1}^D \frac{1}{N} \sum_{n=1}^N f(\mathbf{x}^{(n)}) \partial_\lambda \log p(x_d^{(n)}). \quad (3.6)$$

For notational conciseness, we drop the subscript on n_d and simply use n to identify single samples. Now we compare Equation (3.5) and Equation (3.6) term-wise for $N \rightarrow \infty$

$$\sum_{n=1}^N \mathbb{E}_{X_d \sim p(X_d | \mathbf{x}_{< d}^{(n)})} [f(\mathbf{x}_{< d}^{(n)}, X_d, \mathbf{x}_{> d}^{(n)}) \partial_\lambda \log p(X_d | \mathbf{x}_{< d}^{(n)})] = \sum_{n=1}^N f(\mathbf{x}_{< d}^{(n)}, x_d^{(n)}, \mathbf{x}_{> d}^{(n)}) \partial_\lambda \log p(x_d^{(n)}).$$

In the equation above, we see that SCatR takes the expected value for X_d (left-hand side) and computes it exactly using an explicit sum over the space $\Omega(X_d)$, whereas REINFORCE (right-hand side) uses sampled values. This means, in turn, that the left-hand side is a Rao-Blackwellised version of the right-hand side. Doing this for every d gives us a Rao-Blackwellised version for REINFORCE, i.e., the SCatR estimator. \square

Corollary 3.4 (Bias and Variance). The SCateR estimator is unbiased and its variance is upper-bounded by the variance of REINFORCE.

Proof. This follows immediately from Proposition 3.3, the law of total expectation and the law of total variance Blackwell [1947], Radhakrishna Rao [1945]. \square

Computational complexity. Consider Equation (3.5) and observe that none of the random variables has a sample space larger than $K = \max_d(|\Omega(X_d)|)$. Computing our gradient estimate requires performing three nested sums with lower bound 1 and upper bounds equal to D , K and N , respectively. These summations result in a time complexity of $\mathcal{O}(D \cdot K \cdot N)$. Leveraging the parallel implementation of prefix sums on GPUs Ladner and Fischer [1980], a time complexity of $\mathcal{O}(\log D + \log K + \log N)$ can be obtained, which allows for the deployment of SCateR in modern deep architectures.

Function evaluations. The main limitation of SCateR is the number of function evaluations it requires to provide its estimates. Indeed, each term in the three nested sums of Equation (3.5) involves a different function evaluation, meaning an overall number of at most $D \cdot \max_d(|\Omega(X_d)|) \cdot N$ function evaluations is necessary for every estimate. Fortunately, these evaluations can often be parallelised in modern deep architectures, leading to a positive trade-off between function evaluations and performance in our experiments (Section 6).

Connection to Local Expectation Gradients (LEG). LEG Titsias, Michalis and Lázaro-Gredilla, Miguel [2015] was proposed as a gradient estimator tailored for variational proposal distributions. Similar to the CatLog-Derivative trick, it also singles out variables from a target expectation by summing over the variables' domains. However, it does so one variable at a time by depending on the Markov blanket of those variables. This dependency has the disadvantage of inducing a weighted expression, as variables downstream of a singled-out variable X_i are sampled independently of the values assigned to X_i . In this sense, LEG can intuitively be related to Gibbs sampling where a variable is marginalised given possible assignments to its Markov blanket instead of being sampled.

In contrast, the CatLog-Derivative trick is applicable to any probability distribution with a known factorisation, even in the case where different factors share parameters. It sums out variables by following the variable ordering of this factorisation rather than depending on the Markov blanket. Following the topology induced by the distribution's factorisation naturally avoids the disconnect between summed-out and downstream variables for which LEG required a weighting term and makes the CatLog-Derivative trick more related to ancestral sampling. Moreover, computing these LEG weights adds additional computational complexity. A more technical and formal comparison between LEG and the CatLog-Derivative trick can be found in the appendix.

4 The IndeCateR Gradient Estimator

Using the CatLog-Derivative trick derived in the previous section we are now going to study a prominent special case of multivariate categorical distributions. That is, we will assume that our probability distribution admits the independent factorisation $p(\mathbf{X}) = \prod_{d=1}^D p_d(X_d)$. Note that all D different distributions still depend on the same set of learnable parameters Λ . Furthermore, we subscript the individual distributions p_d as they can no longer be distinguished by their conditioning sets. Plugging in this factorisation into Theorem 3.1 gives us the *Independent Categorical REINFORCE* estimator, or *IndeCateR* for short.

Proposition 4.1 (IndeCateR). Let $p(\mathbf{X})$ be a multivariate categorical probability distribution that depends on a set of parameters Λ and factorises as $p(\mathbf{X}) = \prod_{d=1}^D p_d(X_d)$, then the gradient of a finite expectation $\mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})]$ can be estimated with

$$\sum_{d=1}^D \sum_{x_\delta \in \Omega(X_d)} \partial_\lambda p_d(x_\delta) \frac{1}{N} \sum_{n_d=1}^N f(\mathbf{x}_{\neq d}^{(n_d)}, x_\delta), \quad (4.1)$$

where $\mathbf{x}_{\neq d}^{(n_d)}$ are samples drawn from $p(\mathbf{X}_{\neq d})$.

Proof. We start by looking at the expression in Equation (3.1). Using the fact that we have a set of independent random variables, we can simplify $p(x_\delta | \mathbf{X}_{<d})$ to $p_d(x_\delta)$. As a result, the gradient of the expected value can be rewritten as

$$\partial_\lambda \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X})] = \sum_{d=1}^D \sum_{x_\delta \in \Omega(X_d)} \mathbb{E}_{\mathbf{X}_{<d} \sim p(\mathbf{X}_{<d})} [\partial_\lambda p_d(x_\delta) \mathbb{E}_{\mathbf{X}_{>d} \sim p(\mathbf{X}_{>d})} [f(\mathbf{X}_{\neq d}, x_\delta)]] \quad (4.2)$$

$$= \sum_{d=1}^D \sum_{x_\delta \in \Omega(X_d)} \partial_\lambda p_d(x_\delta) \mathbb{E}_{\mathbf{X}_{\neq d} \sim p(\mathbf{X}_{\neq d})} [f(\mathbf{X}_{\neq d}, x_\delta)] \quad (4.3)$$

Drawing N samples for the $D - 1$ independent random variables $\mathbf{X}_{\neq d}$ and for each term in the sum over d then gives us the estimate stated in the proposition. \square

Example 4.2 (Independent Factorisation). Let us consider a multivariate distribution involving three 3-ary independent categorical random variables. Concretely, this gives us

$$p(X_1, X_2, X_3) = p_1(X_1)p_2(X_2)p_3(X_3), \quad (4.4)$$

where X_1, X_2 and X_3 can take values from the set $\Omega(X) = \{1, 2, 3\}$. Taking this specific distribution and plugging it into Equation (4.1) for the IndeCateR gradient estimator now gives us

$$\sum_{d=1}^3 \sum_{x_\delta \in \{1, 2, 3\}} \partial_\lambda p_d(x_\delta) \frac{1}{N} \sum_{n_d=1}^N f(\mathbf{x}_{\neq i}^{(n_d)}, x_\delta). \quad (4.5)$$

In order to understand the difference between the Log-Derivative trick and the CatLog-Derivative trick, we are going to look at the term for $d = 2$ and consider the single-sample estimate

$$\partial_\lambda p_2(1)f(x_1, 1, x_3) + \partial_\lambda p_2(2)f(x_1, 2, x_3) + \partial_\lambda p_2(3)f(x_1, 3, x_3), \quad (4.6)$$

where x_1 and x_3 are sampled values for the random variables X_1 and X_3 . These samples could be different ones for $d \neq 2$. The corresponding single sample estimate using REINFORCE instead of IndeCateR would be $\partial_\lambda p_2(x_2)f(x_1, x_2, x_3)$.

We see that for REINFORCE we sample all the variables whereas for IndeCateR we perform the explicit sum for each of the random variables in turn and only sample the remaining variables.

In Section 6, we demonstrate that performing the explicit sum, as shown in Example 4.2, leads to practical low-variance gradient estimates. Note how, in the case of $D = 1$, Equation (4.1) reduces to the exact gradient. With this in mind, we can interpret IndeCateR as computing exact gradients for each single random variable X_d with respect to an approximation of the function $\mathbb{E}_{\mathbf{X}_{\neq d} \sim p(\mathbf{X}_{\neq d})} [f(\mathbf{X}_{\neq d}, X_d)]$.

LEG in the independent case. The gap between the LEG and CatLog-Derivative trick gradient estimates grows smaller in the case the underlying distributions factorises into independent factors. The additional weighting function for LEG collapses to 1, resulting in similar expressions. However, two main differences remain. First, LEG does not support factors with shared parameters, although this problem can be mitigated by applying the chain rule. Second, and more importantly, LEG does not allow for different samples to be drawn for different variables. While this difference seems small, we will provide clear evidence that it can have a significant impact on performance (Section 6.4).

5 Related Work

Apart from LEG, another closely related work is the RAM estimator [Tokui and Sato, 2017]. RAM starts by reparametrising the probability distribution using the Gumbel-Max trick followed by a marginalisation. We show in Section 3 that this reparametrisation step is unnecessary, allowing SCateR to explicitly retain the conditional dependency structure. Here we also see the different objectives of our work and Tokui and Sato's. While they proposed RAM in order to theoretically study the quality of control variate techniques, we are interested in building practical estimators that can exploit the structure of specific distributions.

Moreover, just like LEG, Tokui and Sato [2017] did not study the setting of a shared parameter space Λ between distributions, although this being the most common setting in modern deep-discrete and neural-symbolic architectures. Hence, an experimental evaluation for this rather common setting is missing in Tokui and Sato [2017]. By providing such an evaluation, we show that efficiently implemented Rao-Blackwellised gradient estimators for categorical random variables are a viable option in practice when compared to variance reduction schemes based on control variates.

Such variance reduction methods for REINFORCE aim to reduce the variance by subtracting a mean-zero term, called the baseline, from the estimate [Bengio et al., 2013, Paisley et al., 2012]. Progress in this area is mainly driven by multi-sample, i.e., sample-dependent, baselines [Grathwohl et al., 2018, Gregor et al., 2014, Gu et al., 2015, Mnih and Gregor, 2014, Ranganath et al., 2014, Titsias, Michalis and Shi, Jiaxin, 2022, Tucker et al., 2017] and leave-one-out baselines [Kool et al., 2019, 2020, Mnih and Rezende, 2016, Richter et al., 2020, Salimans and Knowles, 2014, Shi et al., 2022]. Variance can further be reduced by coupling multiple samples and exploiting their dependencies [Dimitriev and Zhou, 2021a,b, Dong et al., 2020, 2021, Yin and Zhou, 2019, Yin et al., 2019, Yin et al., 2020]. These latter methods usually introduce bias, which has to be resolved by, for instance, an importance weighing step [Dong et al., 2020].

A general drawback of baseline variance reduction methods is that they often involve a certain amount of computational overhead, usually in the form of learning optimal parameters or computing statistics. This computational overhead can be justified by assuming that the function in the expectation is costly to evaluate. In such cases, it might pay off to perform extra computations if that means the expensive function is evaluated fewer times.

Another popular, yet diametrically opposite, approach to low-variance gradient estimates for categorical random variables is the concrete distribution [Maddison et al., 2017]. The concrete distribution is a continuous relaxation of the categorical distribution using the Gumbel-Softmax trick [Jang et al., 2017]. Its main drawback is that it results in biased gradient estimates. Even though this bias can be controlled by a temperature parameter, the tuning of this parameter is highly non-trivial in practice.

6 Experiments

In this section we evaluate the performance of IndeCateR on a set of standard benchmarks from the literature. Firstly, two synthetic experiments (Section 6.1 and Section 6.2) will be discussed. In Section 6.3, a discrete variational auto-encoder (DVAE) experiment is optimised on three standard datasets. Lastly, we study the behaviour of IndeCateR on a standard problem taken from the neural-symbolic literature (Section 6.4).

We mainly compare IndeCateR to REINFORCE leave-one-out (RLOO) Kool et al. [2019], Salimans and Knowles [2014] and to the Gumbel-Softmax trick. The former is a strong representative of methods reducing variance for REINFORCE while the latter is a popular and widely adopted choice in deep-discrete architectures [Huijben et al., 2022].

We implemented IndeCateR and RLOO in TensorFlow Abadi et al. [2015], from where we also took the implementation of the concrete distribution.¹ Furthermore, all gradient estimators were JIT compiled. The different methods were benchmarked using discrete GPUs and every method had access to the same computational resources. A detailed account of the experimental setup, including hyperparameters and hardware, is given in the appendix.

In the analysis of the experiments we are interested in four main quantities: 1) the time it takes to perform the gradient estimation, 2) the number of drawn samples, 3) the number of function evaluations and 4) the value of a performance metric such as the ELBO. The number of function calls matters in as far that they might be expensive. Note, however, that evaluating the function f is comparatively cheap and trivially parallelisable in the standard benchmarks.

6.1 Synthetic: Exact Gradient Comparison

For small enough problems the exact gradient for multivariate categorical random variables can be computed via explicit enumeration. Inspired by Niepert et al. [2021], we compare the estimates of

$$\partial_{\theta} \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} \left[\sum_{d=1}^D |X_d - b_d| \right] \quad (6.1)$$

to its exact value. Here, θ are the logits that directly parametrise a categorical distribution $p(\mathbf{X}) = \prod_{d=1}^D p(X_d)$ and b_d denotes an arbitrarily chosen element of $\Omega(X_d)$. We compare the gradient estimates from IndeCateR, REINFORCE, RLOO, and Gumbel-Softmax (GS) by varying the number of distributions D and their cardinality.

In Figure 1 we show the empirical bias and variance for the different estimators. Each estimator was given 1000 samples, while IndeCateR was only given a single one. Hence, IndeCateR has the fewest function evaluations as $D \cdot K$ is smaller than 1000 for each configuration. IndeCateR offers gradient estimates close to the exact ones with orders of magnitude lower variance for all three settings. RLOO exhibits the smallest difference in bias, yet it can not compete in terms of variance. Furthermore, the computation times were of the same order of magnitude for all methods. This is in

¹Our code is publicly available at <https://github.com/ML-KULeuven/catlog>.

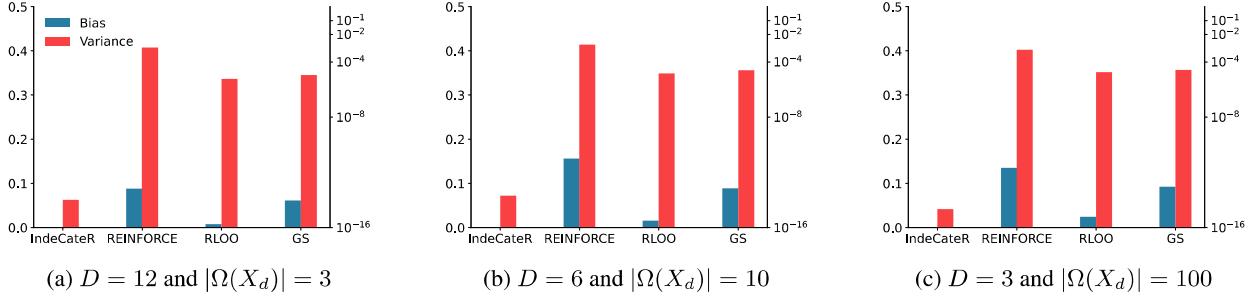


Figure 1: We report the empirical bias and variance for the different estimators and distributions in comparison to the exact gradient. The left y-axis in each plot indicates the bias while the right one shows the variance. Bias and variance results are averages taken over 1000 runs.

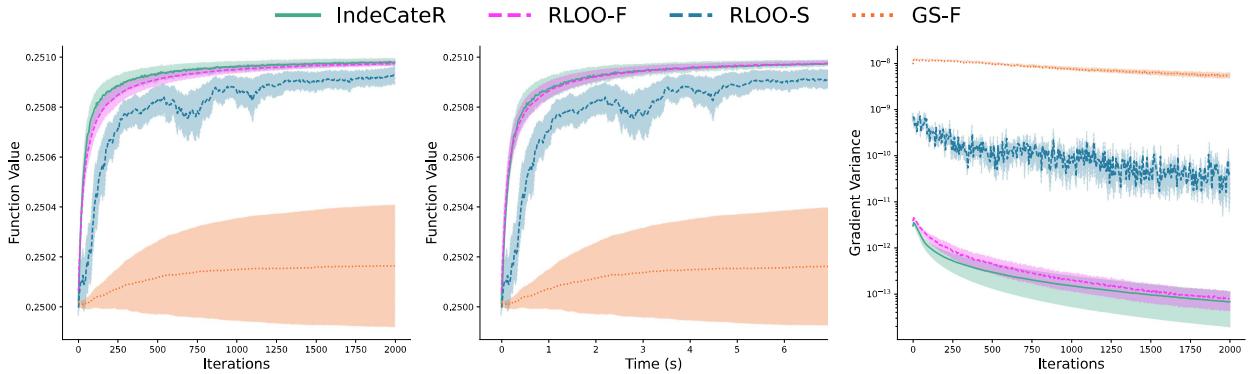


Figure 2: We plot the function value for different estimators against iterations (left) and time (middle). On the right, we plot the variance of the gradients against iterations. Statistics were obtained by taking the average and standard error over 10 runs. IndeCateR and RLOO-S both use 2 samples, while RLOO-F and Gumbel-Softmax (GS) use 800 samples. The number of function evaluations is equal for IndeCateR, RLOO-F, and GS-F. We performed a hyperparameter search for the learning rate and the temperature of GS-F. Parameters were optimised using RMSProp.

stark contrast to the estimator presented by Tokui and Sato [2017], where a two-fold increase in computation time of RAM with respect to REINFORCE is reported.

6.2 Synthetic: Optimisation

We now study an optimisation setting Titsias, Michalis and Shi, Jiaxin [2022], where the goal is to maximise the expected value

$$\mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} \left[\frac{1}{D} \sum_{i=1}^D (X_i - 0.499)^2 \right], \quad (6.2)$$

and $p(\mathbf{X})$ factorizes into D independent binary random variables. The true maximum is given by $p(X_d = 1) = 1$ for all d . This task is challenging because of the small impact of the individual values for each X_d on the expected value for higher values of D . We set $D = 200$ and report the results in Figure 2, where we compare IndeCateR to RLOO and Gumbel-SoftMax.

In Figure 2 and subsequent figures we use the notation RLOO-F and RLOO-S, which we define as follows. If IndeCateR takes N samples, then it performs $D \cdot K \cdot N$ function evaluations with $K = \max_d |\Omega(X_d)|$. As such, we define RLOO-S as drawing the same number of samples as IndeCateR, which translates to N function evaluations. For RLOO-F we match the number of function evaluations, which means that it takes $D \cdot K \cdot N$ samples. We give an analogous meaning to GS-S and GS-F for the Gumbel-SoftMax gradient estimator.

IndeCateR distinguishes itself by having both the lowest variance and quickest convergence across all methods, even compared to RLOO-F. Additionally, the time to compute all gradient estimates does not differ significantly for the

different methods and leads to the same conclusions. It is striking to see that the Gumbel-Softmax estimator struggles in this task, which is likely due to its bias in combination with the insensitive loss function.

6.3 Discrete Variational Auto-Encoder

As a third experiment we analyse the ELBO optimisation behaviour of a discrete variational auto-encoder (DVAE) Rolfe [2017]. We optimise the DVAE on the three main datasets from the literature, being MNIST LeCun [1998], F-MNIST Xiao et al. [2017] and Omniglot Lake et al. [2015]. The encoder component of the network has two dense hidden layers of sizes 384 and 256 ending in a latent 200-dimensional Bernoulli variable. The decoder takes samples from this variable followed by layers of size 256, 384 and 784. IndeCateR again uses two samples, hence we can compare to the same configurations of RLOO and Gumbel-Softmax as in Section 6.2. That is, equal samples (GS-S and RLOO-S) and equal function evaluations (GS-F and RLOO-F).

As evaluation metrics, we show the negated training and test set ELBO in combination with the variance of the gradients throughout training. We opted to report all metrics in terms of computation time (Figure 3), but similar results in terms of iterations are given in the appendix.

A first observation is that IndeCateR performs remarkably well in terms of convergence speed as it beats all other methods on all datasets in terms of training ELBO. However, we can observe a disadvantage of the quick convergence in terms of generalisation performance when looking at the test set ELBO. RLOO-F and IndeCateR both exhibit overfitting on the training data for MNIST and F-MNIST, resulting in an overall higher negative test set ELBO compared to the other methods. We speculate that the relaxation for the Gumbel-Softmax or the higher variance Wu et al. [2020] of RLOO-S act as a regulariser for the network. As one would ideally like to separate discrete sampling layers from network regularisation, we see this overfitting as a feature and not a drawback of IndeCateR.

IndeCateR additionally is competitive in terms of gradient variance, especially considering the number of samples that are drawn. Across all datasets, the gradient variance of IndeCateR is comparable to that of RLOO-F, even though IndeCateR takes only 2 samples compared to the 800 of RLOO-F. The only method with consistently lower gradient variance is the Gumbel-Softmax trick, but only for equal function evaluations. However, in that case, the Gumbel-Softmax trick takes ≈ 1.6 times longer than IndeCateR to compute its gradient estimates.

6.4 Neural-Symbolic Optimisation

A standard experiment in the neural-symbolic literature is the addition of MNIST digits Manhaeve et al. [2018]. We will examine an alternative version of this experiment, that is, given a set of D MNIST digits, predict the sum of those digits. During learning, the only supervision provided is the sum and not any direct label of the digits. The difficulty of the problem scales exponentially, as there are 10^D states in the sample space. Note that there are only $10D + 1$ possible labels, resulting in only very sparse supervision.

In the field of neural-symbolic AI such problems are either solved exactly Manhaeve et al. [2018] or by simplifying the underlying combinatorial structure Huang et al. [2021], Manhaeve et al. [2021]. Exact methods scale very poorly while simplifying the combinatorics introduces problematic biases. In contrast, we will study neural-symbolic inference and learning using sampling and unbiased gradient estimators. The overall architecture is as follows, imitating inference in a general neural-symbolic system. Each of the D different MNIST images is passed through a neural classifier, which gives probabilities for each class. These probabilities are used to sample a number between 0 and 9 for each image. The numbers are summed up and compared to the label using a binary cross-entropy loss, as logical supervision is either true or false.

The hard part of this problem is the sparse supervision. In neural-symbolic AI, this sparsity is resolved by relying on classic search algorithms. However, these are not amenable to be run on parallel computing hardware, such as GPUs. In turn, training has to either occur entirely on the CPU or a hefty time penalty has to be paid for transferring data from and to the GPU.

Using IndeCateR in a neural-symbolic setting we achieve two things. On the one hand, we use sampling as a stochastic yet unbiased search, replacing the usual symbolic search. On the other, we render this stochastic search differentiable by estimating gradients instead of performing the costly exact computation. To make this work, we exploit a feature of IndeCateR that we have ignored so far.

In Equation (4.1) we can draw different samples for each of the D different terms in the sum $\sum_{d=1}^D$ corresponding to the D different variables. While drawing new samples for every variable does increase the total number of samples, the number of function evaluations remains identical. We indicate the difference by explicitly mentioning which version of

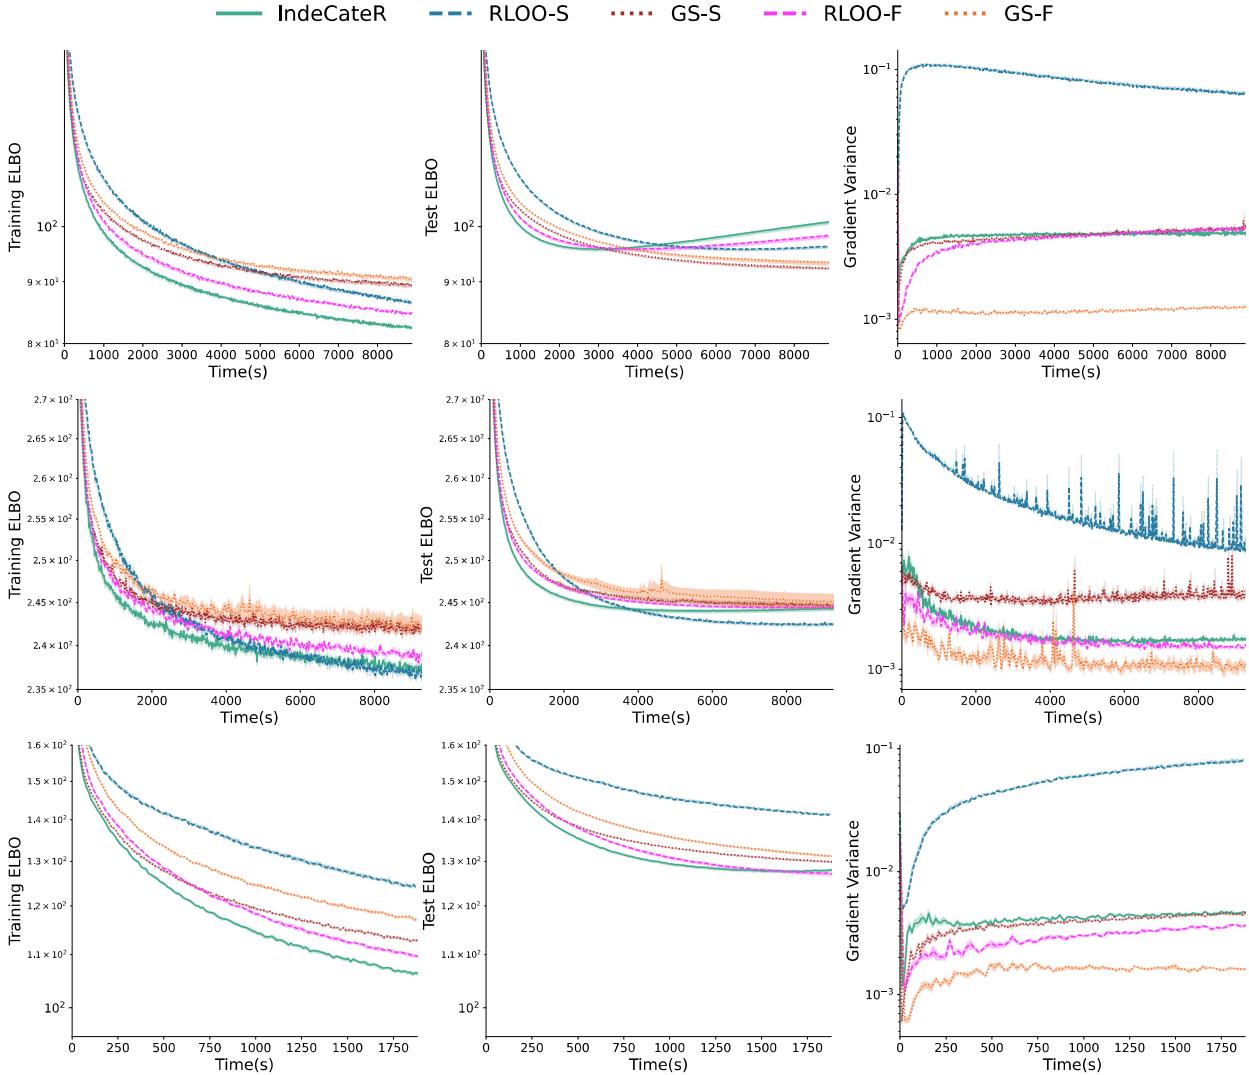


Figure 3: The top row shows negated training ELBO, negated test ELBO and gradient variance for the DVAE on MNIST, each plotted against time. Rows 2 and 3 show the same plots for F-MNIST (middle) and Omniglot (bottom).

IndeCateR draws new samples per variable. Note that this modification of the estimator would not have been possible following the work of Tokui and Sato [2017] or Titsias, Michalis and Lázaro-Gredilla, Miguel [2015].

In Figure 4 we show the empirical results. We let IndeCateR without sampling per variable draw 10 samples such that IndeCateR with sampling per variables draws $D \cdot 10$ samples. The number of function evaluations of both methods stays the same at $D \cdot 10 \cdot 10$. In comparison, RLOO-F uses $D \cdot 10 \cdot 10$ samples and function evaluations. We see that IndeCateR is the only method capable of scaling and consistently solving the MNIST addition for 16 digits when taking new samples for every variable. In particular, these results show how the CatLog-Derivative trick estimates can significantly outperform the LEG and RAM estimates by drawing different samples for different variables, even in the independent case.

7 Conclusion

We derived a Rao-Blackwellisation scheme for the case of gradient estimation with multivariate categorical random variables using the CatLog-Derivative trick. This scheme led straightforwardly to a practical gradient estimator for independent categorical random variables – the IndeCateR estimator. Our experimental evaluation showed that IndeCateR produces faithful gradient estimates, especially when the number of samples is the main concern, and can

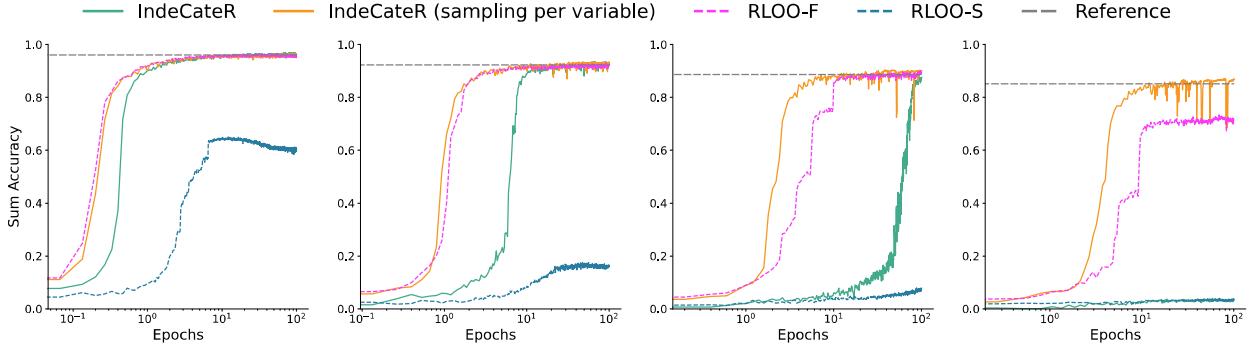


Figure 4: Test set accuracy of predicting the correct sum value versus number of epochs for the MNIST addition. From left to right, the plots show curves for the different estimators for 4, 8, 12, and 16 MNIST digits. We compare IndeCateR with and without sampling per variables to RLOO-F and RLOO-S. As the function f has a zero derivative almost everywhere in this case, we can not compare to the Gumbel-Softmax trick. The dashed, grey line at the top of each plot represents a hypothetical sum-classifier that has access to a 99% accurate MNIST digit classifier. The x -axis is in log-scale.

be run efficiently on modern GPUs. Furthermore, IndeCateR constitutes the only hyperparameter-free, low-variance estimator for (conditionally) independent categorical random variables.

Three main axes can be identified for future work. First and foremost, we aim to analyse the more general case of a non-independently factorising probability distribution by applying SCatE-R to more structured problems, such as differentiable planning and reinforcement learning. Secondly, we plan to further study both IndeCateR and SCatE-R in the neural-symbolic setting and develop estimators for multivariate discrete-continuous distributions containing symbolic structures De Smet et al. [2023]. Finally, the combination of control variates with the CatLog-Derivative trick seems to be fertile ground for further improvements in variance reduction. Considering the Gumbel-Softmax Trick generally still exhibits lower variance for an equal amount of function evaluations, such a combination could put REINFORCE-based methods on equal footing with those based on the Gumbel-Softmax Trick.

Acknowledgements

This research received funding from the Flemish Government (AI Research Program), from the Flanders Research Foundation (FWO) under project G097720N and under EOS project No. 30992574, from the KU Leuven Research Fund (C14/18/062) and TAILOR, a project from the EU Horizon 2020 research and innovation programme under GA No. 952215. It is also supported by the Wallenberg AI, Autonomous Systems and Software Program (WASP) funded by the Knut and Alice Wallenberg- Foundation.

References

- Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dandelion Mané, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. URL <https://www.tensorflow.org/>. Software available from tensorflow.org.
- Yoshua Bengio, Nicholas Léonard, and Aaron Courville. Estimating or propagating gradients through stochastic neurons for conditional computation. *arXiv*, 2013.
- David Blackwell. Conditional expectation and unbiased sequential estimation. *The Annals of Mathematical Statistics*, 1947.
- George Casella and Christian P Robert. Rao-blackwellisation of sampling schemes. *Biometrika*, 1996.
- Lennert De Smet, Pedro Zuidberg Dos Martires, Robin Manhaeve, Giuseppe Marra, Angelika Kimmig, and Luc De Readt. Neural probabilistic logic programming in discrete-continuous domains. *UAI*, 2023.

- Alek Dimitriev and Mingyuan Zhou. Carms: Categorical-antithetic-reinforce multi-sample gradient estimator. *NeurIPS*, 2021a.
- Aleksandar Dimitriev and Mingyuan Zhou. Arms: Antithetic-reinforce-multi-sample gradient for binary variables. *ICML*, 2021b.
- Zhe Dong, Andriy Mnih, and George Tucker. Disarm: An antithetic gradient estimator for binary latent variables. *NeurIPS*, 33, 2020.
- Zhe Dong, Andriy Mnih, and George Tucker. Coupled gradient estimators for discrete latent variables. *NeurIPS*, 2021.
- Artur d'Avila Garcez, Sebastian Bader, Howard Bowman, Luis C Lamb, Leo de Penning, BV Illuminoo, Hoifung Poon, and COPPE Gerson Zaverucha. Neural-symbolic learning and reasoning: a survey and interpretation. *Neuro-Symbolic Artificial Intelligence: The State of the Art*, 2022.
- Will Grathwohl, Dami Choi, Yuhuai Wu, Geoffrey Roeder, and David Duvenaud. Backpropagation through the void: Optimizing control variates for black-box gradient estimation. *ICLR*, 2018.
- Karol Gregor, Ivo Danihelka, Andriy Mnih, Charles Blundell, and Daan Wierstra. Deep autoregressive networks. *ICML*, 2014.
- Shixiang Gu, Sergey Levine, Ilya Sutskever, and Andriy Mnih. Muprop: Unbiased backpropagation for stochastic neural networks. *arXiv*, 2015.
- Jiani Huang, Ziyang Li, Binghong Chen, Karan Samel, Mayur Naik, Le Song, and Xujie Si. Scallop: From probabilistic deductive databases to scalable differentiable reasoning. *NeurIPS*, 2021.
- Iris AM Huijben, Wouter Kool, Max B Paulus, and Ruud JG Van Sloun. A review of the gumbel-max trick and its extensions for discrete stochasticity in machine learning. *IEEE TPAMI*, 2022.
- Eric Jang, Shixiang Gu, and Ben Poole. Categorical reparameterization with gumbel-softmax. *ICLR*, 2017.
- Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *ICLR*, 2015.
- Wouter Kool, Herke van Hoof, and Max Welling. Buy 4 reinforce samples, get a baseline for free! *ICLR Deep RL Meets Structured Prediction Workshop*, 2019.
- Wouter Kool, Herke van Hoof, and Max Welling. Estimating gradients for discrete random variables by sampling without replacement. *ICLR*, 2020.
- Richard E Ladner and Michael J Fischer. Parallel prefix computation. *Journal of the ACM (JACM)*, 27(4):831–838, 1980.
- Brenden M Lake, Ruslan Salakhutdinov, and Joshua B Tenenbaum. Human-level concept learning through probabilistic program induction. *Science*, 350(6266):1332–1338, 2015.
- Yann LeCun. The mnist database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>, 1998.
- Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 1998.
- Chen Liang, Mohammad Norouzi, Jonathan Berant, Quoc V Le, and Ni Lao. Memory augmented policy optimization for program synthesis and semantic parsing. *NeurIPS*, 2018.
- Runjing Liu, Jeffrey Regier, Nilesh Tripuraneni, Michael Jordan, and Jon McAuliffe. Rao-blackwellized stochastic gradients for discrete distributions. *ICML*, 2019.
- Guy Lorberbom, Andreea Gane, Tommi Jaakkola, and Tamir Hazan. Direct optimization through argmax for discrete variational auto-encoder. *NeurIPS*, 2019.
- Guy Lorberbom, Chris J Maddison, Nicolas Heess, Tamir Hazan, and Daniel Tarlow. Direct policy gradients: Direct optimization of policies in discrete action spaces. *NeurIPS*, 2020.
- Chris J Maddison, Andriy Mnih, and Yee Whye Teh. The concrete distribution: A continuous relaxation of discrete random variables. *ICLR*, 2017.
- Robin Manhaeve, Sebastijan Dumancic, Angelika Kimmig, Thomas Demeester, and Luc De Raedt. Deepprobabilog: Neural probabilistic logic programming. *advances in neural information processing systems*, 31, 2018.
- Robin Manhaeve, Giuseppe Marra, and Luc De Raedt. Approximate inference for neural probabilistic logic programming. *KR*, 2021.
- Andriy Mnih and Karol Gregor. Neural variational inference and learning in belief networks. *ICML*, 2014.
- Andriy Mnih and Danilo Rezende. Variational inference for monte carlo objectives. *International Conference on Machine Learning*, 2016.

- Shakir Mohamed, Mihaela Rosca, Michael Figurnov, and Andriy Mnih. Monte carlo gradient estimation in machine learning. *JMLR*, 2020.
- Mathias Niepert, Pasquale Minervini, and Luca Franceschi. Implicit mle: backpropagating through discrete exponential family distributions. *NeurIPS*, 2021.
- John Paisley, David M Blei, and Michael I Jordan. Variational bayesian inference with stochastic search. *ICML*, 2012.
- Max B Paulus, Chris J Maddison, and Andreas Krause. Rao-blackwellizing the straight-through gumbel-softmax gradient estimator. *ICLR*, 2021.
- C Radhakrishna Rao. Information and accuracy attainable in the estimation of statistical parameters. *Bulletin of the Calcutta Mathematical Society*, 1945.
- Rajesh Ranganath, Sean Gerrish, and David Blei. Black box variational inference. *AISTATS*, pages 814–822, 2014.
- Lorenz Richter, Ayman Boustati, Nikolas Nüsken, Francisco Ruiz, and Omer Deniz Akyildiz. Vargrad: a low-variance gradient estimator for variational inference. *NeurIPS*, 2020.
- Jason Tyler Rolfe. Discrete variational autoencoders. *ICLR*, 2017.
- Francisco J. R. Ruiz, Michalis K. Titsias, and David M. Blei. The generalized reparameterization gradient. *NeurIPS*, 2016.
- Tim Salimans and David A Knowles. On using control variates with stochastic approximation for variational bayes and its connection to stochastic linear regression. *arXiv*, 2014.
- Jiaxin Shi, Yuhao Zhou, Jessica Hwang, Michalis Titsias, and Lester Mackey. Gradient estimation with discrete stein operators. *NeurIPS*, 2022.
- Titsias, Michalis and Lázaro-Gredilla, Miguel. Local Expectation Gradients for Black Box Variational Inference. *NeurIPS*, 2015.
- Titsias, Michalis and Shi, Jiaxin. Double control variates for gradient estimation in discrete latent variable models. *AISTATS*, 2022.
- Seiya Tokui and Issei Sato. Evaluating the variance of likelihood-ratio gradient estimators. *ICML*, 2017.
- George Tucker, Andriy Mnih, Chris J Maddison, John Lawson, and Jascha Sohl-Dickstein. Rebar: Low-variance, unbiased gradient estimates for discrete latent variable models. *NeurIPS*, 2017.
- Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Reinforcement learning*, 1992.
- Jingfeng Wu, Wenqing Hu, Haoyi Xiong, Jun Huan, Vladimir Braverman, and Zhanxing Zhu. On the noisy gradient descent that generalizes as sgd. In *ICML*. PMLR, 2020.
- Han Xiao, Kashif Rasul, and Roland Vollgraf. Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*, 2017.
- Mingzhang Yin and Mingyuan Zhou. Arm: Augment-reinforce-merge gradient for stochastic binary networks. *ICLR*, 2019.
- Mingzhang Yin, Yuguang Yue, and Mingyuan Zhou. Arsm: Augment-reinforce-swap-merge estimator for gradient backpropagation through categorical variables. *ICML*, 2019.
- Mingzhang Yin, Nhat Ho, Bowei Yan, Xiaoning Qian, and Mingyuan Zhou. Probabilistic best subset selection via gradient-based optimization. *arXiv*, 2020.

A Experimental Details

We will give more details about our experimental setup in this section. We used two machines for all experiments. The first one has access to 4x RTX 3080 Ti coupled with an Intel Xeon Gold 6230R CPU @ 2.10GHz and 256 GB of RAM. The second one has 4x GTX 1080 Ti with an Intel Xeon CPU E5-2630 v4 @ 2.20GHz and 128 GB of RAM. All synthetic experiments and the MNIST and F-MNIST datasets for the DVAE experiment were run on the former, while Omniglot and the neural-symbolic experiment used the latter machine. Note that every single run of each experiment only used a single GPU.

A.1 Synthetic Experiments

Modelling. Both synthetic experiments directly modelled their respective distributions via parametrised logits. That is, the parameters are given by a $D \times K$ matrix of logits, where D and K are the number of independent components and the dimension of each individual variable, respectively.

Hyperparameters. We optimised the temperature value for the Gumbel-Softmax for both synthetic experiments and the learning rate for all methods for the second one via a grid search over $\{10, 1, 0.1, 0.01\}$. This yielded a temperature value of 1 for the first synthetic experiment and of 0.1 for the second one. In addition to this choice of initial temperature, a standard exponential annealing scheme Titsias, Michalis and Shi, Jiaxin [2022] was also used during the synthetic optimisation experiment, multiplying the initial temperature value with a factor $e^{-0.05}$ every 20 iterations. Moreover, the Gumbel-Softmax gradients also required removal of diverging values to maintain their usability during optimisation. Finally, the learning rates for all methods were chosen from the values $\{10, 5, 1, 0.1, 0.01\}$ while using the RMSProp optimiser and ended up being different for almost all methods. Both RLOO-F and IndeCateR managed to perform best for a higher learning rate of 5 while RLOO-S was stable and performed best for a learning rate of 1. Gumbel-Softmax needed the smallest learning rate of 0.01 as it would otherwise exhibit extremely unstable behaviour, likely due to the sensitivity of the loss function in conjunction with its bias. The choice of encoder architecture indeed has 2 true hidden layers and one output layer. The choice of two samples for IndeCateR was made as it is the lowest number of samples that allows a sample-equivalent to RLOO as it is a multi-sample estimator.

A.2 Discrete Variational Auto-Encoder

Datasets. Three datasets were used in this experiment with different forms of pre-processing. First, the usual MNIST LeCun [1998] dataset was binarised, as per usual in the discrete gradient estimation literature Niepert et al. [2021], Titsias, Michalis and Shi, Jiaxin [2022]. Binarised means that all pixels are normalised to values between 0 and 1 by dividing by 255 followed by replacing all values strictly higher than 0.5 with one and those lower with 0. For F-MNIST Xiao et al. [2017], we kept the continuous form of the data, only normalising the pixels to values between 0 and 1. In this way, we also had a dataset with a continuous supervision signal. The last dataset, Omniglot Lake et al. [2015], was both downscaled and binarised. Every Omniglot image usually has dimensions $105 \times 105 \times 3$, which we reduced to greyscale 28×28 followed by a binarisation step as outlined earlier. All test sets are constructed in the same way from the dedicated test partitions of the respective datasets.

Modelling. All datasets used the exact same neural architecture for the discrete variational auto-encoder (DVAE). As the data all had the same input size of 28×28 , the first layer flattened the input to size 784 followed by three dense layers of sizes 384, 256 and 200. The first two of these layers used the ReLU activation function while the last one had no non-linearity. The output of this final 200-dimensional layer predicted the logits of a 200-dimensional Bernoulli variable, of which samples are drawn. These samples then formed the input to the decoder component of the DVAE, consisting of another three dense layers of size 256, 384 and 784. Again, the first two layers utilised ReLU activation functions while the final layer had a linear activation. Finally, following the literature, the output of the decoder is interpreted as the logits for 784 binary random variables and optimised using an ELBO loss function, which is an expected value. The correct probabilities are given by the normalised and binarised pixel values of the original image.

Hyperparameters. As in the synthetic experiment, the Gumbel-Softmax required the most optimisation. The initial temperature value was set at 1 and annealed using an exponential decay with a factor of $e^{-0.01}$ applied at the start of every epoch. To guarantee stability, it was necessary to both limit the temperature to 0.1 in conjunction with removing diverging values from the gradients. This effort emphasises how the Gumbel-Softmax requires a lot of additional care to make it functional in practice compared to methods like RLOO or IndeCateR. Apart from the optimisation of the Gumbel-Softmax, we tried a learning rate of both 10^{-3} and 10^{-4} for all methods and ended up picking the latter because it generally provided lower final negated ELBO values due to its slower convergence. All methods used the

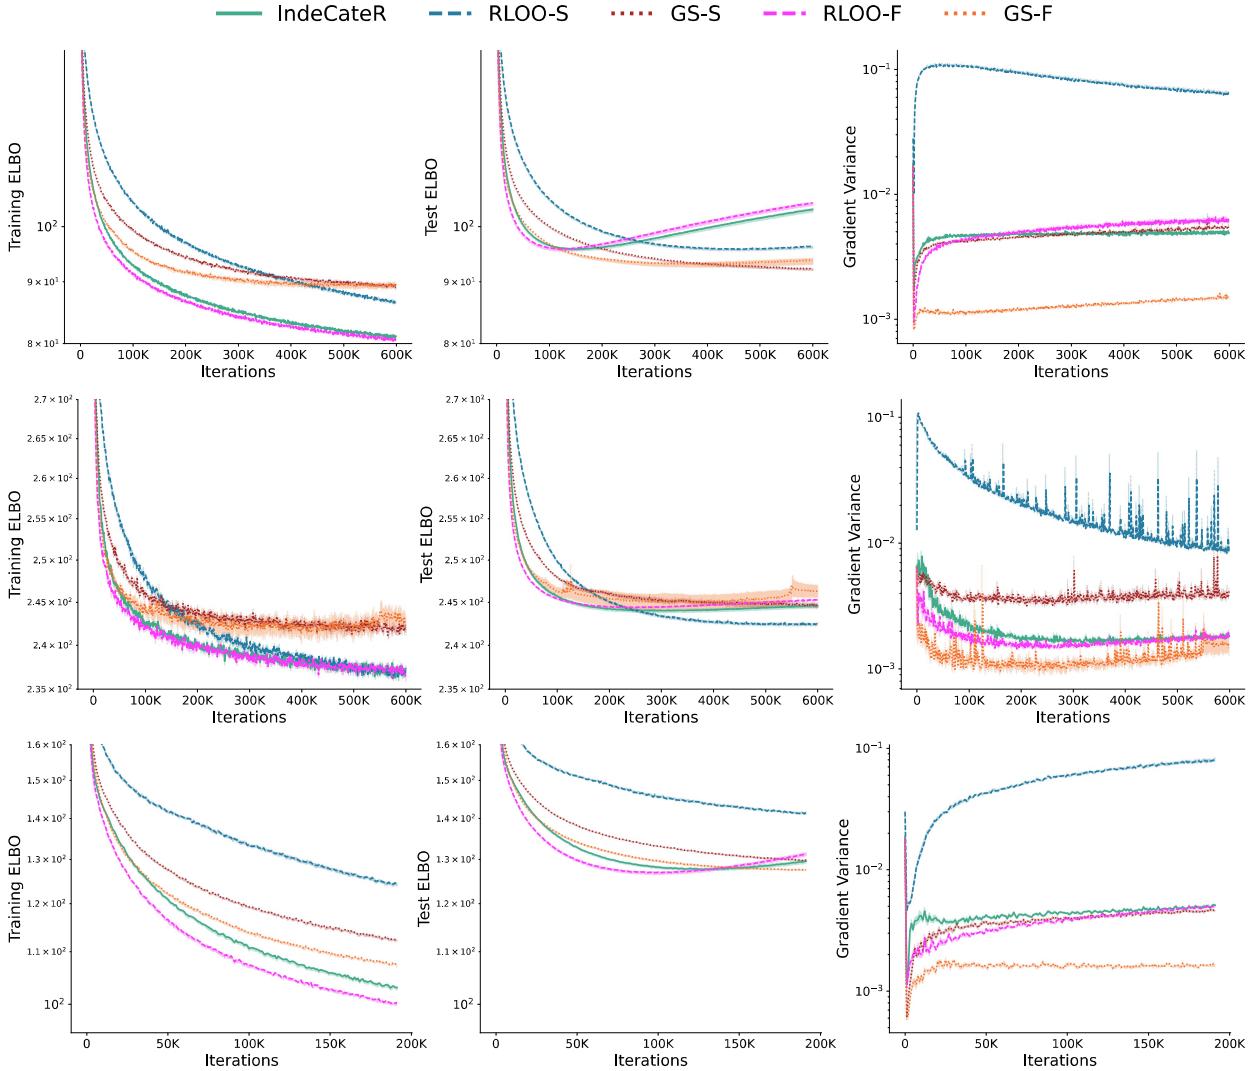


Figure 5: The top row shows negated training ELBO, test ELBO and gradient variance for the DVAE on MNIST – each plotted against iterations. Rows 2 and 3 show the same plots for F-MNIST (middle) and Omniglot (bottom). Omniglot only shows around 200K iterations because the dataset is around one third the size of MNIST and F-MNIST, and we fix the number of epochs to 1000.

Adam Kingma and Ba [2015] optimiser. Note that all networks were initialised using the same random scheme and that no weight regularisation was applied to avoid obscuring the impact of using different estimators.

Additional plots and interpretations. We report additional results in Figure 5 for the DVAE, where we plot all metrics in terms of iterations. Using these figures, we can see that IndeCateR is very close to the performance of RLOO-F, but with orders of magnitude less samples (2 for IndeCateR compared to 800 for RLOO-F). This reduction in samples is also what makes IndeCateR perform better in time. The same overfitting behaviour can now also be seen on Omniglot for both IndeCateR and RLOO-F. However, in this case, the lowest negated test ELBO for both methods is lower than all competitors. This indicates that adding separate regularisation might further improve performance, as this regularisation is implicitly present in GS-S, GS-F and RLOO-S through either the continuous relaxations or higher variance.

A.3 Neural-Symbolic Optimisation

Dataset. The data for performing the MNIST addition experiment Manhaeve et al. [2018] on sequences of N digits is easily generated by randomly selecting N images from the MNIST dataset and concatenating them into a sequence. The supervision of such a sequence is equally easily obtained by summing up the labels of the sampled MNIST images. We only allow every MNIST image to appear once in all sequences, meaning the dataset contains $\lfloor 60000/N \rfloor$ sequences to learn from. A similar procedure holds for the test set, whose sequences we construct from the test set partition of the MNIST dataset.

Modelling. The neural network used for this experiment is a traditional LeNet LeCun et al. [1998] consisting of two convolutional layers with 6 and 16 filters of size 5 with ReLu activations followed by a flattening operation. Next, there are three dense layers of sizes 120, 84 and 10 of which the first two also have ReLU activations and the final one activates linearly. The output of this network for every image in the sequence yields the logits for the independent categorical distributions from which we sample. These samples are summed up and supervised to the correct sum. That is, a sampled sum is supervised to 1 if it matches the correct sum or 0 otherwise. More specifically, we optimise the negative log likelihood

$$-\log P\left(\sum_i^N D_i = s\right) = -\log \mathbb{E}_{D \sim p(D)} \left[\mathbb{1}_{\sum_i^N D_i = s} \right]. \quad (\text{A.1})$$

Hyperparameters. We chose the standard Adam optimiser with a learning rate of 10^{-3} because of its dependable performance. The choice of 10 samples for IndeCateR was made so that every possible outcome of each digit is expected to appear once when sampling with uniform probabilities. The number of samples for the other methods is then fixed.

A.4 Brittleness of Gumbel-Softmax Estimates

During our experimental process, a number of noteworthy difficulties arose when applying the Gumbel-Softmax trick (GS). These difficulties show how applying the GS is far from plug-and-play and requires a substantial amount of hyperparameter tuning, at least to achieve performance comparable to other methods like RLOO.

Temperature tuning. While the bias of the GS can theoretically be controlled by annealing the temperature to zero, this process turns out to be rather cumbersome in practice. For example, in the DVAE experiments, we had to fix a lowest value for the temperature value beyond which it could not be annealed. If not, the ever smaller temperature would lead to numerical instability in the form of diverging gradients. Even when lower limiting the temperature, sporadic divergences of the gradient still needed to be clipped or they would instantly disrupt the full training process. In other words, finding a well-performing annealing scheme is non-trivial and does not guarantee a problem-free training process.

Bias and insensitive objectives. As observed during the synthetic optimisation experiment, the loss function itself can seemingly have a significant impact on the quality of GS gradient estimates. This experiment used a regular mean squared error loss function, which does not have any non-differentiabilities in its domain. However, the value of the loss function itself does not change much for different sampled values. Since the GS estimates are biased, we conjecture that there are cases where the bias overpowers the information from the loss function. In theory, annealing the temperature should eventually make the bias a non-issue, but the numerical issues raised in the previous paragraph prevented us from acquiring a satisfactory solution.

B Extended Related Work

There are a plethora of works in the literature of gradient estimators, for an in-depth discussion we refer the reader to the survey in Mohamed et al. [2020]. This section will extend the related work given in the main paper.

Continuous relaxations. Diametrically opposite to REINFORCE-based methods are those based on continuous relaxations of categorical distributions. The concrete distribution Maddison et al. [2017] is a continuous relaxation of the categorical distribution using the Gumbel-Softmax trick Jang et al. [2017]. Compared to variance reduction techniques based on control variates for REINFORCE, the Gumbel-Softmax trick has the advantage of being a single-sample estimator. However, it has other drawbacks. Firstly, it gives us a biased estimate of the gradients. Secondly, it necessitates the tuning of the temperature hyperparameter – either through hyperparameter search or a specific annealing

schedule Jang et al. [2017], Maddison et al. [2017]. Thirdly, when the derivative of the function inside the expectation is zero almost everywhere, such as the reward in reinforcement learning, the Gumbel-Softmax trick always returns zero gradients. Even though techniques have been developed to improve upon some of these shortcomings, such as Rao-Blackwellising the concrete distribution Paulus et al. [2021], the obtained improvements seem marginal. For a more detailed account of the concrete distribution and related methods we refer the reader to Huijben et al. [2022].

Coupling multiple samples. Coupling approaches make better use of multiple samples by making them dependent. Initially, these approaches were applied to the case of a multivariate binary distribution Dimitriev and Zhou [2021b] by, for example, using a logistic reparametrisation together with antithetic sampling Yin and Zhou [2019], Dong et al. [2020], Yin et al. [2020]. Categorical extensions of these binary cases have also been proposed. One line of work Dimitriev and Zhou [2021a] extends the binary version by representing a categorical distribution as a series of binary ones. Another Yin et al. [2019] utilises a Dirichlet reparametrisation together with different kinds of coupled sampling. Our Cat-Log Derivative trick is also orthogonal to coupling methods, as it currently only considers independent samples.

Control variate methods. While some control variate techniques were mentioned in the main paper, a couple of more intricate contributions need to be mentioned. RELAX [Grathwohl et al., 2018] and REBAR [Tucker et al., 2017] are hybrid methods that use continuous relaxations to construct a control variate for REINFORCE. The work of Titsias, Michalis and Shi, Jiaxin [2022] introduces double control variates; a global one for all samples and a local one for each sample separately. Our Cat-Log Derivative trick is orthogonal to control variate approaches in general and therefore both strategies could benefit from each other.

Rao-Blackwellisation. Apart from the RAM and LEG gradient estimators, there are two more lines of work that apply Rao-Blackwellisation. The works of Lorberbom et al. [2019, 2020] use a finite difference approximation to compute the gradient. Other approaches Liu et al. [2019], Liang et al. [2018] globally sum categories while sampling from the remainder. In contrast, our approach conditions the sampling at the level of the variables and can be therefore again be regarded as a complementary strategy.

C Comparing CatLog and Local Expectation Gradients

Local Expectation Gradients (LEG) Titsias, Michalis and Lázaro-Gredilla, Miguel [2015] and the CatLog-Derivative trick both try to decompose the overall gradient of an expectation by exploiting distributional knowledge. However, there are a couple of fundamental theoretical differences, which we collect in the following statement.

Proposition C.1. (LEG Trick and Estimator) Given a factorised distribution $p(\mathbf{X}) = \prod_{d=1}^D p(X_d | \mathbf{X}_{<d})$, LEG with respect to a shared parameter $\lambda \in \Lambda$ leads to the expression

$$\begin{aligned} & \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X}) \partial_\lambda \log p(\mathbf{X})] \\ &= \sum_{d=1}^D \mathbb{E}_{\mathbf{X}_{\neq d} \sim p(\mathbf{X}_{\neq d})} [\mathbb{E}_{X_d \sim p(X_d | \text{MB}(X_d))} [f(\mathbf{X}) \partial_\lambda \log p(X_d | \mathbf{X}_{<d})]] \\ &\quad (\text{cf. Eq. 8 in Titsias, Michalis and Lázaro-Gredilla, Miguel [2015]}) \\ &= \sum_{d=1}^D \mathbb{E}_{\mathbf{X}_{<d} \sim p(\mathbf{X}_{<d})} \left[\mathbb{E}_{X'_d \sim p(X'_d | \mathbf{X}_{<d})} \left[\mathbb{E}_{\mathbf{X}_{>d} \sim p(\mathbf{X}_{>d} | X'_d, \mathbf{X}_{<d})} \left[\mathbb{E}_{X_d \sim p(X_d | \mathbf{X}_{\neq d})} [f(\mathbf{X}) \partial_\lambda \log p(X_d | \mathbf{X}_{<d})] \right] \right] \right] \end{aligned} \tag{C.1}$$

where $\text{MB}(X_d)$ is the Markov blanket of variable X_d and

$$p(X_d | \text{MB}(X_d)) = p(X_d | \mathbf{X}_{\neq d}) = \frac{p(X_d | \mathbf{X}_{<d}) \prod_{j>d} p(X_j | \mathbf{X}_{<j})}{\sum_{x_\delta \in \Omega(X_d)} p(x_\delta | \mathbf{X}_{<d}) \prod_{j>d} p(X_j | \mathbf{X}_{<j})} \tag{C.2}$$

is a weighting distribution. The Monte Carlo estimate corresponding to the LEG Trick in Equation (C.1) (cf. Algorithm 1 in Titsias, Michalis and Lázaro-Gredilla, Miguel [2015]) given N pivot samples is given by the following expression:

$$\mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X}) \partial_\lambda \log p(\mathbf{X})] \approx \frac{1}{N} \sum_{n=1}^N \sum_{d=1}^D \sum_{x_\delta \in \Omega(X_d)} p(x_\delta | \mathbf{x}_{\neq d}^{(n)}) f(\mathbf{x}_{\neq d}^{(n)}, x_\delta) \partial_\lambda \log p(x_\delta | \mathbf{x}_{<d}^{(n)})$$

Additionally, the computational complexity is $\mathcal{O}(N \cdot D^2 \cdot K)$, with $K = \max_d(|\Omega(X_d)|)$.

Proof. We start by recalling the proof for the first equality in Equation (C.1).

$$\begin{aligned}
\mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X}) \partial_\lambda \log p(\mathbf{X})] &= \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} \left[f(\mathbf{X}) \partial_\lambda \log \prod_{d=1}^D p(X_d | \mathbf{X}_{<d}) \right] \\
&= \sum_{d=1}^D \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [f(\mathbf{X}) \partial_\lambda \log p(X_d | \mathbf{X}_{<d})] \\
&= \sum_{d=1}^D \mathbb{E}_{\mathbf{X}_{\neq d} \sim p(\mathbf{X}_{\neq d})} \left[\underbrace{\mathbb{E}_{X_d \sim p(X_d | \mathbf{X}_{\neq d})} [f(\mathbf{X}) \partial_\lambda \log p(X_d | \mathbf{X}_{<d})]}_{= g(\mathbf{X}_{\neq d})} \right]. \tag{C.3}
\end{aligned}$$

By noting that $p(X_d | \mathbf{X}_{\neq d}) = p(X_d | \text{MB}(X_d))$, we obtain the first result.

Regarding the second equality in Equation (C.1), we observe that

$$\mathbb{E}_{\mathbf{X}_{\neq d} \sim p(\mathbf{X}_{\neq d})} [g(\mathbf{X}_{\neq d})] \tag{C.4}$$

in Equation (C.3) can be equivalently rewritten as

$$\begin{aligned}
\mathbb{E}_{\mathbf{X}_{\neq d} \sim p(\mathbf{X}_{\neq d})} [g(\mathbf{X}_{\neq d})] &= \mathbb{E}_{\mathbf{X} \sim p(\mathbf{X})} [g(\mathbf{X}_{\neq d})] \\
&= \mathbb{E}_{\mathbf{X}_{<d} \sim p(\mathbf{X}_{<d})} \mathbb{E}_{X'_d \sim p(X'_d | \mathbf{X}_{<d})} \mathbb{E}_{\mathbf{X}_{>d} \sim p(\mathbf{X}_{>d} | X'_d, \mathbf{X}_{<d})} [g(\mathbf{X}_{\neq d})] \tag{C.5}
\end{aligned}$$

By plugging this result into Equation (C.3) we obtain the desired result for the second equality in Equation (C.1).

The Monte Carlo estimate trivially follows by sampling N pivots from the first three expectations in Equation (C.1) and by taking the exact weighted average of score evaluations.

Regarding computational complexity, we observe that computing the weighting function in the Monte Carlo estimate requires $\mathcal{O}(D)$ evaluations. Also, the estimate requires to iterate over three summations, thus having a computational requirement of $\mathcal{O}(N \cdot D \cdot K)$. By combining these two results, we obtain the overall complexity of $\mathcal{O}(N \cdot D^2 \cdot K)$. \square

Note that this result does not immediately follow from the expressions provided by Titsias, Michalis and Lázaro-Gredilla, Miguel [2015], as they do not explicitly study the case of a shared parameter space across different variables. We do show, however, that their assumption of a separate parameter space can be removed, allowing a one-to-one comparison to the CatLog-Derivative trick. We give an overview of the main differences between LEG and the CatLog-Derivative trick in Table 1. To be precise, we will refer to the following two expressions for LEG and CatLog that clearly indicate the differences between the two.

$$\text{LEG} = \sum_{d=1}^D \mathbb{E}_{(\mathbf{x}_{<d}, \textcolor{red}{X}'_d, \mathbf{x}_{>d}) \sim p(\mathbf{x}_{<d}, \textcolor{red}{X}'_d, \mathbf{x}_{>d})} [\mathbb{E}_{\mathbf{x}_d \sim p(\mathbf{x}_d | \mathbf{x}_{\neq d})} [f(\mathbf{X}) \partial_\lambda \log p(X_d | \mathbf{X}_{<d})]] \tag{C.6}$$

$$\text{CatLog} = \sum_{d=1}^D \mathbb{E}_{(\mathbf{x}_{<d}, \textcolor{red}{X}_d, \mathbf{x}_{>d}) \sim p(\mathbf{x}_{<d}, \textcolor{red}{X}_d, \mathbf{x}_{>d})} [f(\mathbf{X}_{\neq d}, X_d) \partial_\lambda \log p(X_d | \mathbf{X}_{<d})] \tag{C.7}$$

While the most important theoretical difference is the presence of the weighing factor $p(X_d | \text{MB}(X_d)) = p(X_d | \mathbf{X}_{\neq d})$, there is also a practical difference. Indeed, IndeCateR allows for new samples to be drawn for each variable while Algorithm 1 in Titsias, Michalis and Lázaro-Gredilla, Miguel [2015] does not do so. This difference translates into a significant differential in performance as observed in the neural-symbolic experiment.

Table 1: Main differences between LEG and CatLog. The different nature of the applied tricks result in different computational complexities and related sampling strategies.

Name	Trick	Complexity	Relation to Sampling
LEG	Equation (C.6)	$\mathcal{O}(N \cdot D^2 \cdot K)$	Gibbs sampling
CatLog	Equation (C.7)	$\mathcal{O}(N \cdot D \cdot K)$	Ancestral sampling